

Administrator Guide — Kubernetes Auto Resource Cleanup (Deployments, Pods & Services)

Audience: Cluster Administrators (operations / platform owners). This guide explains how the Auto Cleanup tool works, how to deploy and operate it safely, and what administrator actions are required when incidents occur.

1. Purpose and Scope

The Auto Cleanup tool automatically frees stale resources in the Kubernetes cluster to preserve compute (GPU/CPU), reduce contention, and keep the DGX environment usable for all users. It enforces time-based policies (soft and hard limits) with administrative controls for exclusions and logging.

This document explains the **operational** steps an administrator needs to take — installation, configuration, scheduling, monitoring, exclusions, and incident responses.

2. Key Concepts

- **Hard limit** — strict maximum age (minutes). When a resource reaches this age it is deleted immediately.
 - **Soft limit** — earlier age (minutes). When a resource reaches this age it becomes a candidate for deletion; the presence of a `keep-alive: "true"` label prevents deletion under soft-limit conditions.
 - **Keep-alive label** — a user-level label (`metadata.labels.keep-alive`) that prevents soft-limit deletes when set to `true`.
 - **Exclusions** — administrator-maintained lists that exempt whole namespaces or individual resources from deletion.
 - **Batch pod deletion** — the script collects pods that should be deleted and issues batched `kubectl delete pod pod1 pod2 ...` commands so deletions proceed in parallel and the script does not block.
 - **Locking** — only one instance of the script runs at a time to prevent cron overlap.
 - **Logging** — every action (decision and delete) is logged with timestamps to a centralized log file.
-

3. Files delivered with the tool

Place these files in the same directory on your control node (head node):

- `auto-pod-deletion.sh` — the main script (executable)
- `cleanup_config.env` — runtime configuration
- `exclude_namespaces.txt` — namespaces to skip (one per line)
- `exclude_deployments.txt` — specific deployment names to skip

- `exclude_pods.txt` — specific pod names to skip
- `exclude_services.txt` — specific service names to skip
- `README.md` — user-facing overview (already supplied)

Keep the directory owned by root and readable only by administrators.

4. Prerequisites & environment assumptions

- The script runs as a user with permission to delete Deployments, Pods, and Services in the `dgx-*` namespaces.
 - A persistent location for logs (e.g. `/var/log/giindia/auto_cleanup_logs/`). Ensure the `LOG_FILE` path in `cleanup_config.env` is writable by the script user.
 - Cron or another scheduler for periodic execution.
-

5. Installation and initial setup (administrator steps)

1. Upload the script and config files to the control node (use an admin account):

```
scp auto-pod-deletion.sh cleanup_config.env admin@headnode:/opt/auto-cleanup/
cd ~/auto-cleanup/
sudo chown root:root *
sudo chmod 755 auto-pod-deletion.sh
```

1. Edit `cleanup_config.env` to match your policy (examples and defaults are included in the file). Key variables to set are `STUDENT_SOFT`, `STUDENT_HARD`, `FACULTY_SOFT`, `FACULTY_HARD`, `INDUSTRY_SOFT`, `INDUSTRY_HARD`, and `LOG_FILE`.

2. Create the log directory and ensure permissions:

```
sudo mkdir -p /var/log/giindia/auto_cleanup_logs
sudo chown root:root /var/log/giindia/auto_cleanup_logs
sudo chmod 755 /var/log/giindia/auto_cleanup_logs
```

1. (Optional) Create the exclusion files if you need to protect specific namespaces or resources:

2. `exclude_namespaces.txt` — one namespace per line (example `dgx-admin`)
3. `exclude_deployments.txt` — one deployment name per line
4. `exclude_pods.txt` — one pod name per line
5. `exclude_services.txt` — one service name per line

6. Dry-run validation (recommended):

7. Temporarily set `Pod=False` and `Service=False` in `cleanup_config.env` and run the script to validate logging and scanning only.

8. Check logs for expected scanning behavior. Once validated, enable desired resources and re-run.

6. Configuration details (what each key means)

This section is a quick reference for policy-maintenance.

General toggles

- `Deployment`, `Pod`, `Service` — set `True` or `False`. When `False`, that entire resource type is skipped.

Limit feature toggles

- `Deployment_HardLimit`, `Deployment_SoftLimit` — control whether the hard/soft path is active for deployments.
- `Pod_HardLimit`, `Pod_SoftLimit` — control pod behaviors.
- `Service_HardLimit`, `Service_SoftLimit` — control service behaviors.

If both hard and soft are set to `False` for a resource, the script treats that resource category as disabled.

Time limits (minutes)

- `STUDENT_SOFT`, `STUDENT_HARD` — minutes applied to namespaces matching `dgx-s*`.
- `FACULTY_SOFT`, `FACULTY_HARD` — minutes applied to `dgx-f*`.
- `INDUSTRY_SOFT`, `INDUSTRY_HARD` — minutes applied to `dgx-i*`.

Set values in minutes. Example: `STUDENT_SOFT=120` (2 hours).

Pod deletion performance controls

- `POD_BG_DELETE` (or similar flag in the deployed script) — background deletion toggle. When true the script issues deletes in background and moves on.
- `POD_FORCE_DELETE` — if set, the script uses `--grace-period=0 --force` when deleting pods (use with caution).
- `POD_BATCH_SIZE` — number of pods deleted per `kubectl delete` batch (tune for API server limits).

The script comes configured with safe defaults. Only change batch size when you experience API throttling or want faster throughput.

7. Running the script via cron (recommended production setup)

Edit root crontab with `sudo crontab -e` and add a scheduled job:

```
# Run the cleanup at minute 0 every hour (example)
0 * * * * /bin/bash /opt/auto-cleanup/auto-pod-deletion-final-v3.sh >> /var/
log/giindia/auto_cleanup_cron.log 2>&1
```

Guidelines:

- Keep the schedule longer than the worst-case runtime of the script (default run is short because pods are deleted in background). An hourly schedule is typical.
- Point the cron stderr/stdout to a cron-specific log for quick checks.

8. Monitoring & logs

- **Primary Log:** the `LOG_FILE` path from `cleanup_config.env` — contains timestamps and actions.
- **Cron Log:** `/var/log/giindia/auto_cleanup_cron.log` — captures cron stdout/stderr.

What to look for:

- Large numbers of actions in a single run (mass deletions) — indicates an abnormal cleanup event.
- Repeated errors from `kubectl` — might indicate cluster connectivity, RBAC, or permission issues.
- Unprocessed pods stuck in `Terminating` — check node status and finalizers on pods.

Sample troubleshooting commands:

```
# tail last 200 lines of the main log
sudo tail -n 200 /var/log/giindia/auto_cleanup_logs/auto_cleanup.logs

# check for Terminating pods
kubectl get pods -A | grep Terminating

# describe a stuck pod
kubectl describe pod <name> -n <namespace>
```

9. Incident handling & admin actions

A. Accidental deletion (common scenarios)

If a resource was deleted and must be recovered:

1. Confirm deletion time and resource identifier from logs.
2. Check if a backup (user-supplied or stored) exists. If none, the resource must be recreated.
3. If multiple deletions happened due to misconfiguration, revert `cleanup_config.env` immediately and re-run the script in Dry-Run mode (set `Pod=False` etc.) to identify remaining candidates.

Note: The administrator cannot reliably recover data from deleted pods — users should store important output in persistent volumes, object storage, or Git.

B. Too many deletions in a single run

1. Immediately disable the offending resource type: update `cleanup_config.env` (set `Pod=False` or `Deployment=False`) and re-run the script once.
2. Inspect logs to identify which policy (soft/hard) triggered deletion.
3. Notify affected users and restore from backups if available.

C. Pods stuck in Terminating

Common causes: finalizers, volume attachments, node unreachable. Steps:

1. `kubectl describe pod <pod> -n <ns>` to inspect finalizers and events.
2. If finalizers prevent deletion and you must force removal, use
`kubectl patch pod <pod> -n <ns> -p '{"metadata":{"finalizers":null}}' --type=merge` (use sparingly).
3. Investigate node status: `kubectl get nodes` and `kubectl describe node <node>`.

10. Security & RBAC

- The account running the script must have restricted RBAC permissions: `get, list, delete` on Deployments, Pods, and Services only in the target namespaces.
- Do **not** run the script with cluster-admin unless absolutely necessary.
- Protect `cleanup_config.env` and exclusion files with strict permissions (owner=admin, mode=600 or 640).

Recommended minimal RBAC policy (example Role):

```
kind: Role
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: cleanup-role
  namespace: <target-namespace> # or use multiple roles per namespace
rules:
- apiGroups: ["apps"]
  resources: ["deployments"]
  verbs: ["get", "list", "delete"]
- apiGroups: [""]
  resources: ["pods", "services"]
  verbs: ["get", "list", "delete"]
```

Bind the Role to a ServiceAccount used by the control node if you prefer not to use an admin user.

11. Maintenance & best practices

- Review logs daily (or weekly) during initial rollout.
 - Communicate to users the required `keep-alive` label and soft/hard limits.
 - Maintain a recovery SOP for users (how to persist outputs to PVC / object storage).
 - Tune `POD_BATCH_SIZE` if you receive throttling from the API server. Start with 20-50.
-

12. Appendix — Quick reference commands

```
# Dry run: disable deletes in config and run script
# tail logs
sudo tail -f /var/log/giindia/auto_cleanup_logs/auto_cleanup.logs

# Force delete a specific pod immediately
kubectl delete pod <pod> -n <namespace> --grace-period=0 --force

# Show pods stuck in Terminating
kubectl get pods -A | grep Terminating

# Remove finalizers (use with caution)
kubectl patch pod <pod> -n <ns> -p '{"metadata":{"finalizers":null}}' --
type=merge
```

13. Document history

- v1.0 — Initial admin guide (aligned with Auto Cleanup v3).
-

Prepared for cluster administrators. For developer-oriented design & source please reference the README and implementation scripts.